


This article was downloaded by: [Universitat Jaume I]

On: 11 July 2014, At: 00:27

Publisher: Taylor & Francis

Informa Ltd Registered in England and Wales Registered Number: 1072954 Registered office: Mortimer House, 37-41 Mortimer Street, London W1T 3JH, UK



International Journal of Geographical Information Science

Publication details, including instructions for authors and subscription information:

<http://www.tandfonline.com/loi/tgis20>

Geosocial gauge: a system prototype for knowledge discovery from social media

Arie Croitoru^a, Andrew Crooks^b, Jacek Radzikowski^a & Anthony Stefanidis^a

^a Geography and Geoinformation Science, George Mason University, Fairfax, VA, USA

^b Computational Social Science, George Mason University, Fairfax, VA, USA

Published online: 06 Sep 2013.

To cite this article: Arie Croitoru, Andrew Crooks, Jacek Radzikowski & Anthony Stefanidis (2013) Geosocial gauge: a system prototype for knowledge discovery from social media, International Journal of Geographical Information Science, 27:12, 2483-2508, DOI: [10.1080/13658816.2013.825724](https://doi.org/10.1080/13658816.2013.825724)

To link to this article: <http://dx.doi.org/10.1080/13658816.2013.825724>

PLEASE SCROLL DOWN FOR ARTICLE

Taylor & Francis makes every effort to ensure the accuracy of all the information (the "Content") contained in the publications on our platform. However, Taylor & Francis, our agents, and our licensors make no representations or warranties whatsoever as to the accuracy, completeness, or suitability for any purpose of the Content. Any opinions and views expressed in this publication are the opinions and views of the authors, and are not the views of or endorsed by Taylor & Francis. The accuracy of the Content should not be relied upon and should be independently verified with primary sources of information. Taylor and Francis shall not be liable for any losses, actions, claims, proceedings, demands, costs, expenses, damages, and other liabilities whatsoever or howsoever caused arising directly or indirectly in connection with, in relation to or arising out of the use of the Content.

This article may be used for research, teaching, and private study purposes. Any substantial or systematic reproduction, redistribution, reselling, loan, sub-licensing, systematic supply, or distribution in any form to anyone is expressly forbidden. Terms &

Geosocial gauge: a system prototype for knowledge discovery from social media

Arie Croitoru^{a*}, Andrew Crooks^b, Jacek Radzikowski^a and Anthony Stefanidis^a

^a*Geography and Geoinformation Science, George Mason University, Fairfax, VA, USA;*

^b*Computational Social Science, George Mason University, Fairfax, VA, USA*

(Received 7 February 2013; accepted 13 June 2013)

The remarkable success of online social media sites marks a shift in the way people connect and share information. Much of this information now contains some form of geographical content because of the proliferation of location-aware devices, thus fostering the emergence of *geosocial* media – a new type of user-generated geospatial information. Through geosocial media we are able, for the first time, to observe human activities in scales and resolutions that were so far unavailable. Furthermore, the wide spectrum of social media data and service types provides a multitude of perspectives on real-world activities and happenings, thus opening new frontiers in geosocial knowledge discovery. However, gleaning knowledge from geosocial media is a challenging task, as they tend to be unstructured and thematically diverse. To address these challenges, this article presents a system prototype for harvesting, processing, modeling, and integrating heterogeneous social media feeds towards the generation of geosocial knowledge. Our article addresses primarily two key components of this system prototype: a novel data model for heterogeneous social media feeds and a corresponding general system architecture. We present these key components and demonstrate their implementation in our system prototype, GeoSocial Gauge.

Keywords: social media; geosocial analysis; social network analysis; event monitoring; system architecture

1. Introduction

The ubiquity of social media in almost every aspect of modern life is evident in the rapid growth of the sheer volume of data generated by its users. In 2011, Twitter users were posting approximately 200 million tweets per day (Twitter 2011). A year later, this doubled to 400 million (Forbes 2012), reaching a worldwide rate of over 270,000 tweets per minute. At the same time, Flickr users upload in excess of 3000 images per minute (Sapiro 2011), and YouTube users upload approximately 72 hours of video per minute (YouTube 2013). These are remarkable volumes of user-generated data, signifying the shift that has occurred in recent years in digital data production. Although in the past established government or commercial organizations were responsible for generating most of the digital data, today it is estimated that approximately 75% of all digital data is contributed by individual users (Mearian 2011). This trend in data growth is expected to become even more significant over the next several years (Hollis 2011), as computing and technological advances

*Corresponding author. Email: acroitor@gmu.edu

are solidifying the role of the general public as the main contributor and consumer of big data.

Coincident with these trends is the proliferation of location-aware devices. This makes a large portion of user-generated content contributed through web services to be geolocated, thus fostering the emergence of a new type of geospatial information: user-generated, geolocated (or georeferenced) multimedia feeds of diverse thematic content. These feeds are of diverse value, because they are expressions of geo-journalism, conveying current information about significant events, ranging from political movements and uprising (New York Times 2011, Pollock 2011) to natural disasters (Earle *et al.* 2010, Crooks *et al.* 2013). These feeds also communicate users' opinions and views (Bruns and Burgess 2011, Tumasjan *et al.* 2011) or even communicate their experiential perception of the space around them (as in the concept of urban imaginaries of Kelley 2011). As a result, we argue that social media feeds are becoming increasingly *geosocial* in the sense that they often have a substantial geographical content. This can take the form of coordinates from which the contributions originate or of references to specific locations. At the same time, information on the underlying social structure of the user community can be derived by studying the interactions between users (e.g., formed as the respond to, or follow, other users), and this information can provide additional context to the data analysis. Geosocial media therefore emerges as an alternative form of geo-information, which, through its volume and richness, opens new avenues and research challenges for the understanding of dynamic events and situations.

This rise of geosocial media represents a deviation from the well-established concepts of crowdsourcing and volunteered geographic information (VGI Goodchild 2007, Fritz *et al.* 2009, Goodchild and Glennon 2010): unlike Wikimapia or OpenStreetMap, social media feeds are not a vehicle for citizens to explicitly and purposefully contribute geographic information to update or expand geographic databases. Instead, the geographic content is embedded in the contributors' comments and has to be harvested and analyzed before it can be utilized. This type of geographic information has been referred to as Ambient Geographic Information (AGI, Stefanidis *et al.* 2013) and represents an extension of the VGI concept. One could argue that while VGI is primarily crowdsourcing, with specific tasks outsourced to the public at large, AGI is crowd harvesting, with the general public broadcasting information that can be harvested in a meaningful manner.

The utilization of multiple geosocial media sources for information extraction and knowledge generation in various application domains is a challenging task, both in terms of data management and analysis and in terms of knowledge production. For example, many of the current efforts, such as the Ushahidi¹ platform, provide a means to collect data from multiple social media sources and disseminate it over the web. While such efforts are valuable for aggregating and visualizing data, they currently lack capabilities to add context to content or to support detailed analysis. The heterogeneity and diversity of geosocial media, which are both defining properties and essential to their value, often result in a lack of structural homogeneity and adherence to standards that are typically present in more traditional authoritative sources (e.g., government generated data). As a result, the form of raw geosocial media tends to be unstructured, and valuable knowledge is often implicit and cannot be easily processed through automation (Sahito *et al.* 2011). Such data structure heterogeneity has a direct impact on the ability to store it in a single integrated database and manage or process effectively. This heterogeneity is accompanied by thematic diversity: user activity in Twitter can range, for example, from daily chatter and conversations to news reporting and sharing or seeking information (Java *et al.* 2009). Accordingly, the

ability to integrate content across different feeds presents a challenge that requires a holistic approach to the analysis and synthesis of such data.

To address these challenges, this article presents a system prototype for harvesting, processing, modeling, and integrating heterogeneous social media feeds towards the generation of geosocial knowledge. Rather than focusing on a specific geosocial media source or a specific user type, our framework aims to empower users to harvest, ingest, and analyze diverse data sources (starting with Twitter and Flickr). We focus on two key components of this system prototype: a general system architecture for harvesting geosocial data and a novel data model for heterogeneous social media feeds. We present these key components and demonstrate results from their implementation in *GeoSocial Gauge*, our system prototype. Through these results we showcase the unique insights we can glean by integrating spatial and social analysis, especially as it relates to identifying the footprint of distributed connected communities, and observing the effects of various events on these communities and the corresponding social media activity. More specifically, we demonstrate how we find the social structure of these communities, their spatial footprint, and how these communities are affected by relevant events. Together, these elements provide us with a new lens for observing and understanding the human landscape as a geosocial ecosystem, comprising communities, their locations, and the links that connect them. The analysis presented here provides some examples of geosocial knowledge discovery, as defined in the context of this article, and is representative of the capabilities of systems such as *GeoSocial Gauge*. These examples can be viewed as addressing some of the broader challenges identified by Sui and Goodchild (2011) regarding the convergence of GIS and social media.

The remainder of this article is organized as follows. In Section 2, we present our approach to modeling geosocial data, outline the key elements of this model, and discuss the relationships between the different model components. In Section 3, we present *GeoSocial Gauge*, our Geosocial analysis workbench, including its workflow, the data harvesting process, and the design principles of the workspace environment. To demonstrate the benefits of our framework, we describe in Section 4 our analysis results from two recent real-world sociopolitical events. Finally, we outline in Section 5 our conclusion and outlook assessment.

2. A system architecture for geosocial knowledge discovery

In the context of this article, we refer to geosocial knowledge as the aggregate understanding of human activities reflecting the interactions among individuals and groups of people and how these relate to space and time. This represents an extension of traditional geospatial data gathering and processing, expanding its scope beyond physical infrastructure, terrain, and buildings, to include the complex social-cultural fabric of the world, as represented both in the physical and the cyber realms.

A key characteristic of social media is that much of the information that can be derived from it is implicit and has to be extracted through analysis. For example, while it is often possible to derive a social network from social media feeds, this typically requires additional analysis to make user connections explicit. Harvesting geosocial knowledge therefore requires developing a dedicated system that would be able to collect social media feeds, analyze them to reveal any implicit information, and store this information in a database to support knowledge discovery.

Based upon these observations, we have designed a dedicated system built around two guiding principles:

- As social media data delivery is highly dynamic, the system is designed to ingest feeds at a high rate independently of any further analysis that is carried out. The decoupling of the initial data collection from further analysis steps eliminates bottlenecks in the workflow and provides a higher degree of fault tolerance.
- The analysis of feeds is performed using a multi-step approach – ranging from preliminary analysis that involves only basic verification to in-depth analysis in which data are broken into atomic building blocks, and implicit links between data elements are identified and created.

As the events communicated in social media are often dynamic and evolving, their harvesting requires a *continuous* refinement process. This process starts by collecting data from one social media source, using, for example, specific regions of interest, time intervals, and keywords. The harvested social media data are then analyzed, and the results are used to refine the data gathering process, for example, by refining the search area, time intervals, or keywords used for data collection. In doing so, social media is not regarded simply as an additional data source, but rather as a key driver for the data collection process. As the rate of contributions in social media varies significantly among platforms, our data harvesting refinement process starts with the collection of data from a highly dynamic source, that is Twitter, and as the harvesting process is refined, other less dynamic social media sources (e.g., Flickr or YouTube) are harvested. This allows enriching the data available for further analysis, thus enhancing our ability to glean knowledge from it.

2.1. Harvesting and ingesting geosocial data

The transformation of social media feeds into geosocial data is a multi-step process that generally requires three sequential steps: data collection from a set of social media data providers via Application Programming Interfaces (APIs), processing its geolocation content, and storing the data in a dedicated database for further analysis. There exist a number of tools that perform parts of these processes, such as 140 kit,² or hootsuite,³ but these are limited in their scalability with respect to large data sets. However, currently available tools offer limited capabilities to add context to content or to support detailed analysis. In conjunction, there are some noteworthy efforts to create systems for gleaning situational awareness from Twitter and other web text documents (e.g., Tomaszewski *et al.* 2011, MacEachren *et al.* 2011a, 2011b), but they lack the broader scope we pursue here. Consequently, we have designed and developed a dedicated system for harvesting and ingesting geosocial media feeds. As we discuss below, a key characteristic of this system is the ability to continuously harvest, process, and store social media feeds while minimizing any dependencies between these processes. This design minimizes the effect of processing bottlenecks on the overall system performance. Figure 1 below depicts a schematic view of our system. The harvesting process depicted in this figure is driven by Twitter content and then expanded to incorporate Flickr and YouTube; however, any of these social media sources may be used to drive the process. Nevertheless, even though from a database standpoint the order of these services can be interchanged, the nature of these services suggests that Twitter should indeed lead, as shown in Figure 1, owing to its highly dynamic nature and its immediacy in reporting events (e.g., Blasingame 2011, Crooks *et al.* 2013).

Generally, feeds can be retrieved from social media services using their APIs. This allows user-created client programs to interface with the service, submit queries in the form of an HTTP request, and receive in response data in XML or JSON formats. The query parameters may be, for example, based on location (area of interest to which the

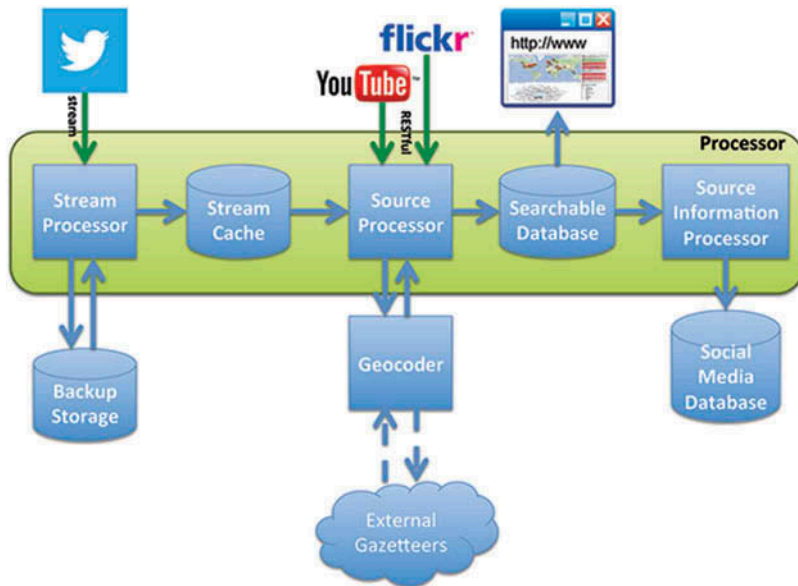


Figure 1. Architecture of a system to harvest and ingest social media feeds that is driven by Twitter.

feed is related), time (period of interest), content elements (keywords), or even user handle or unique ID. It should be noted that depending on the service, some queries may require client authentication (e.g., a user-specific API key), while other queries can be submitted without authentication. Considering Twitter in particular, two types of APIs can be used. The first is streaming, where Twitter content is fed to the client continuously, subject to the used filter parameters (e.g., keywords, location) and volume limits. The second API type is representational state transfer (RESTful, Fielding 2000), where the client issues queries to the server at discrete time instances and Twitter returns the corresponding results. In response to issued queries and depending on the particularities of the social media source that is being queried, either only metadata or data and metadata are returned.

When it comes to Flickr, the information extracted from the RESTfulAPI comprises exclusively metadata (e.g., author, time, and geolocation when available) and information on how to access the actual image itself. In contrast, the Twitter API provides both the actual tweets and the related metadata. Accordingly, from a system perspective, a key difference between the two is that the Flickr API queries do not result in the data itself, but only in metadata including information necessary for retrieving the actual data (i.e., images).

In our system, a query is submitted once a connection with the API is established, and the stream processor receives the returned responses. The stream processor then checks that the data are valid and calculates basic data volume statistics and saves the data to a storage backup in case streaming rates are higher than processing rates. Once this initial processing is completed, the data are transferred to the stream cache. Implemented as a collection in mongoDB, the stream cache is used to isolate the stream processor from the source processor. This allows the two processors to work independently, protecting the collection process from the impact of potential problems and delays in the processing of Tweets.

Data from the stream cache are transferred to the source processor, which has three primary functions with respect to Twitter: parsing the information from the tweets that are stored as JSON streams, geocoding, and the insertion of processed data into the source database. Of these three processes, geocoding is the most critical, as it is responsible for making location data explicit, thus enhancing the value of social media data for geosocial analysis. In our system, geocoding is performed by a dedicated geocoder component using one or more external gazetteers and is activated when explicit geocoding (i.e., coordinates) is not available in a data item. In addition, the source processor is responsible for retrieving and processing data from RESTful APIs (e.g., Flickr and YouTube).

After the geocoding process, data from the source processor are transferred to the searchable database. This database stores the processed social media data as individual data elements, and like the stream cache, it is also implemented using MongoDB to support fast access for near real-time basic analysis. For example, data from the social media database can be used for visualization in mash-ups that require quick access to individual tweets that were subject to the first level processing. The goal of such visualizations is to provide timely dissemination and exploration of the data for fast discovery of trends or the identification of anomalies, which does not require full knowledge of the underlying structure of the connections (i.e., the social network) between different social media users.

The final component of our system involves processing the data and storing it in a dedicated database to support further analysis to reveal connections between social media users and requires finer querying capabilities. A key differentiator between the source information processor and the previous two processors is that it supports social network analysis by making the connections between users explicit. This processing step is separated from the previous two processing steps as it is slower, requiring multiple transactions with a dedicated database. The outputs of the source information processor are stored in the social media database, which is implemented as an SQL database using PostgreSQL. This is a searchable repository of data that supports integration of the various data structures under a unified scheme, thus allowing performing multi-source geosocial analysis. We discuss in Section 3 this unified scheme.

3. Modeling heterogeneous geosocial data

The analysis and utilization of heterogeneous social media requires the development of a conceptual data model that will allow the integration of the various data structures under a unified scheme. Generally, this task can be viewed as a data-cleaning problem, that is, the removal of errors and inconsistencies in databases (Rahm and Do 2000), from either a single source or multiple sources of data. For multiple (heterogeneous) data sources, data-cleaning problems can arise at the schema level (e.g., structural inconsistencies) or at the instance level (e.g., uniqueness violations). In this work, we assume consistency within each source (e.g., Twitter or Flickr) as each social media site maintains its own database. However, since we are interested in integrating multiple social media data sources, structural inconsistencies should be considered. For example, different sources may have certain elements that are unique or specific to them. This problem has been recently indicated in the context of processing social network data in business applications (Bonchi *et al.* 2011) or in the context of music recommender systems (Tan *et al.* 2011).

A step towards a more general solution for integrating social data was recently presented by Lyons and Lessard (2012), who introduced a social feature integration technique for existing information systems that are not socially oriented. However, to the best of our knowledge, there are still no widely accepted models that could be directly applied

to multiple social media sources for geographical analysis. For example, Sahito *et al.* (2011) presented a framework for enriching and deriving linkages in Twitter data by using semantic web resources, such as DBpedia,⁴ FreeBase,⁵ and GeoNames.⁶ However, this work considers only a single social media source – Twitter. A more generalized conceptual model has been recently introduced by Reinhardt *et al.* (2010). In this model, social media data are seen as a combination of two intertwined networks: a set of artifact networks that describes the relationships between data elements (e.g., chats, blogs, or wiki articles) and a social network. This artifact actor network (AAN) model is created by linking the two networks through semantic relationships. Content elements, for example, chats or wiki articles, are modeled using two fundamental building blocks: actors (i.e., users) and artifacts (i.e., content). Although the overall framework concept of the AAN is not tied to a specific social media type, its implementation for different data sources is tailored to each source, for example, Twitter data are modeled differently than wiki sources. Our work is closely related to the data model presented by Shimojo *et al.* (2010), which focuses on lifelogs: digital records of the experiences and events a person encounters during a period of time, which are generated by individuals (Kalnikaitė *et al.* 2010). Their work presents a model that is geared towards the integration of multiple heterogeneous social media sources through the introduction of a common model for lifelog data. Such lifelog data, as captured by online services such as Twitter, Facebook, or Flickr, are modeled according to six fundamental perspectives: (1) what (the data content), (2) why (the purpose of the data), (3) when (the time when the data entry occurred), (4) who (the initiator of the data entry and other users involved), (5) where (the location at which the data entry was made), and (6) how (the means by which the lifelog was recorded). Based on this schema, the modeling process involves mapping data source elements to the relevant common data model perspectives and converting the data to a common format. However, the basic data model unit used in this work is the lifelog, and relations between lifelogs (and users) are not modeled.

The data model presented in this article complements and extends both models presented above and is designed around two fundamental principles:

- (1) Social media has both source-independent and source-dependent components: due to the nature of the social media production process, some components of the data, for example, a user, a data entry, and timestamp, will always be present, while others, such as a photo URL, may be present in some data sources (e.g., Flickr) and not in others (e.g., Twitter).
- (2) Social media is inherently interconnected in a variety of ways reflecting relations among authors and content. These relations may be active, whereby an author, for example, may comment or reply to another contribution. They may also be passive, whereby, for example, an author may simply be following other authors (by subscribing to their feeds).

We argue that this decoupling of source-dependent and source-independent components allows us to reach a high level of abstraction in our model, and diminishes the need to tailor a specific solution for each data source type. By taking advantage of the source-independent component of social media entries we can turn unstructured data streams into structured information, which can be analyzed to extract knowledge. This allows us, for example, to identify social structures that span several social services, or references to the same event or topic coming from different sources and/or services. In Figure 2 we provide an entity–relation diagram of our data model, demonstrating the use of entry, geolocation,

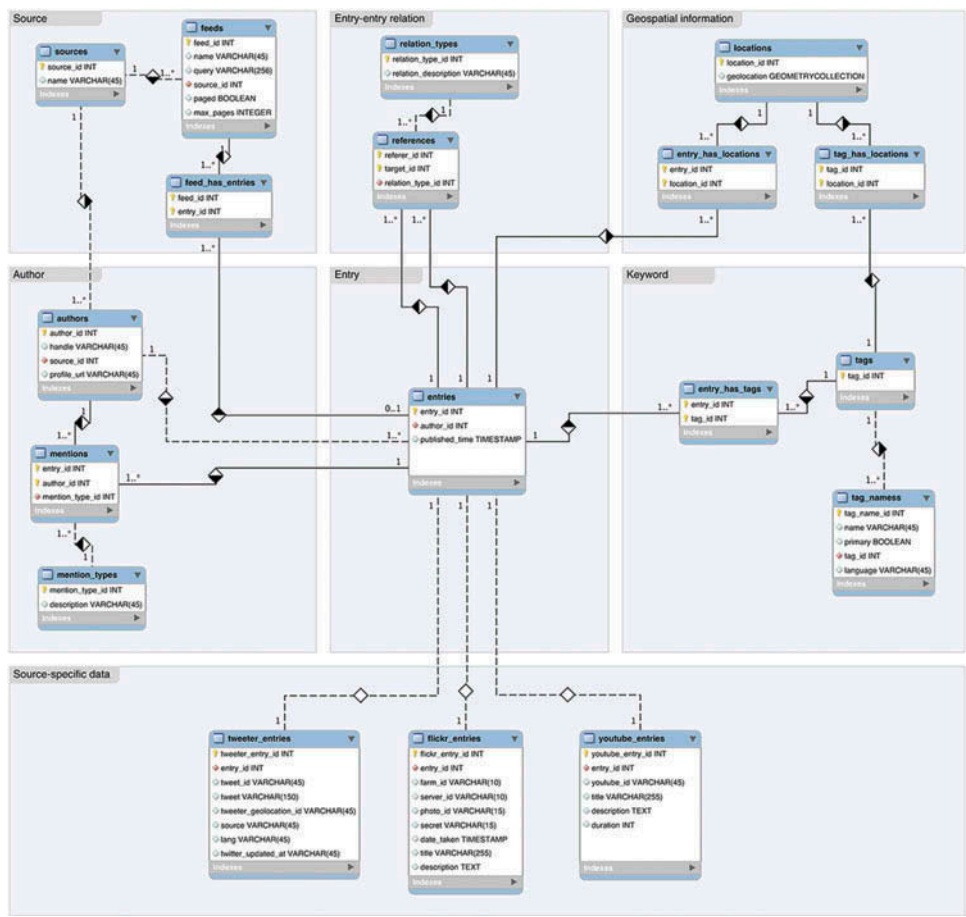


Figure 2. An Entity-Relation (ER) diagram of our geosocial media data model. Related entities are connected with a line, where the cardinality of the relation is denoted by the diamond at the center of the line and the labels at the ends of the line. Relations with a white diamond denote a one-to-one relation, and relations with a white and black diamond denote a one-to-many relation. The cardinality of each side of the relation is denoted by a label at each end of the line (e.g., ‘1 . . . *’ represents one or more and ‘0 . . . 1’ represents zero or one).

time, keywords, and authors to integrate diverse social media feeds. We describe below key components of this diagram.

3.1. Source-independent components

3.1.1. Entry component

An information entry is regarded in our model as an abstract record of social media data, a high level entity that binds the various components that constitute a data entry. Because of their inherent interconnected nature, our model accommodates many-to-many relationships between entries through a references table. In our model, the entry component must be linked to a single source-dependent component and that it must be linked to at least one source feed (e.g., various queries to Twitter, Flickr, YouTube). An entry instance is

uniquely identified by an *entry_id* and is linked to one author instance of the author component. In addition, each entry is associated with a timestamp indicating when was the entry published (for further details on the time component, please refer to Section 3.1.4). Optionally, an entry can have links to other source-independent components. For example, each Twitter entry can be associated with one or more keywords that are represented by records in the tags table.

3.1.2. Author component

In our model, social media contributors (users) are associated with their contributions in a one-to-many relationship. As an author is identified by a tuple of a user name and a social media service identifier, different users can use the same identifier on different services. While some work has addressed the development of strategies to identify the same users across different systems (Iofciu *et al.* 2011), this extends beyond the scope of this article. It should be noted that authors can also be referenced to in the content of social media feeds, for example, a reference to a person's username in a tweet or a response to content submitted by a user. Such references are important as they allow for the establishment of a network of relations between users, which enables the reconstruction of the underlying social network. In our model, this is accomplished by linking entries to users through the mentions table. The type of the reference (e.g., mention of a username, response, etc) is defined by assigning specific attribute from the *mention_type* table.

3.1.3. Geolocation component

Geolocation information for social media feeds can be inferred *indirectly* from content analysis or it can be extracted *directly* from the data itself. Representative examples of the indirect inference of geolocation for social media contributions include Topsy's geo-inference capabilities⁷ and MetaCarta's GeoTag module,⁸ which geolocate contributions based on an analysis of their content and its references to geographic toponyms. While geotagging such contributions is an interesting multidisciplinary problem (see Fink *et al.* 2009, Larson *et al.* 2011) in the sense that it links contributions to locations referenced in them, this is beyond the scope of this work.

Direct geolocation is the method exploited by our system. It is included in the contributions themselves, either in the form of exact coordinates or as a toponym (e.g., listing a city name) to be geolocated using any gazetteer service (e.g., Lieberman *et al.* 2010). Within a tweet record, for example, geolocation information may be populated either by the publisher (Twitter) of the client through which the tweet was submitted, and it may be available in the *place*, *coordinates*, or *location* fields. Similarly, in Flickr, the geolocation information is available in the location record at various levels of granularity, ranging from country name to city, neighborhood, and precise coordinates. It can also be extracted from the image's exchangeable image file format (EXIF) record. A more detailed discussion on the various forms of geolocation has been offered by several authors, including Marcus *et al.* (2011), Croitoru *et al.* (2012), and Crooks *et al.* (2013).

Reports on the percentage of social media information that has geolocation information associated with it vary rather widely. For example, 6% of the approximately one billion Facebook users have elected to enter their actual home address (Backstrom *et al.* 2010). Regarding Twitter, we have reports of tweets carrying geolocation information as toponyms at a city level or better ranging from 20% (Cheng *et al.* 2010) to 65% (Hecht *et al.* 2011) of the total number of tweets. Geolocation information in the form of precise coordinates

is available for a smaller percentage of tweets, with reports ranging from 5% (Cheng *et al.* 2010, Valkanas and Gunopulos 2013) to 16% (Stefanidis *et al.* 2013). This variation in the rate of geolocated tweets can be attributed to a range of factors, such as geographic area, time, and theme. For example, the seemingly high rate of 16%, which was observed following the 2011 tsunami and Fukushima disaster in Japan, can be attributed to the response of the Japanese population to a major event (e.g., evacuation), the increased use of mobile devices after the event (Kaigo 2012), and the high penetration rate of Twitter in Japan (26.6%⁹). A recent study by Leetaru *et al.* (2013) using Twitter's Dekahose data stream has reported that even in the absence of a significant disruptive event (e.g., a nuclear disaster), the rate of geolocated tweets may vary depending on the time of day (from 2.3% at 1:00 pm PST to 1.7% at 6:00 am PST) or the geographic location (from 2.86% in Jakarta to 0.77% in Moscow).

3.1.4. Time component

While geolocated information is available for only a percentage of social media feeds, temporal information is available for all of them. For tweets, this is their submission time (the instance when it was submitted to Twitter by the user). A few seconds (2–5 s) after their submission, the tweets become available to the public through Twitter. With Flickr, we can have two time records: image capture (through its EXIF timestamp) or image posting on Flickr. Experiments have shown that social media response to events may be nearly real-time, with tweets reporting, for example, earthquakes as early as few tens of seconds after them (Earle *et al.* 2010, Earle *et al.* 2011, Crooks *et al.* 2013, Kraut *et al.* 2013).

3.1.5. Narrative component: keywords

Social media are in essence modern reporting and storytelling mechanisms: they enable individuals to communicate their impressions and interests to the general public. As part of this narrative, participants contribute keywords, both explicitly and implicitly. The explicit contribution of keywords is through tags, like hashtags in Twitter (Huang *et al.* 2010) or general tags in Flickr (Ames and Naaman 2007). These keyword tags are used to emphasize content, facilitate data access, and enable community networking (Romero *et al.* 2011). Hashtag usage, for example, has been shown to accelerate data retrieval from Twitter (Zhao *et al.* 2011) and Flickr (Sun *et al.* 2011). Hashtags also support the building of semantic networks by allowing individual tweets to be linked thematically based on their content. All tweets, for example, with the hashtag #obama are linked together at the Twitter site. Similarly, photos sharing the same tags in Flickr are aggregated into photosets (the equivalent of single-user photo albums) and groups (multi-user albums).

In addition to these explicit tagging mechanisms, keyword may also emerge by being adopted massively, for example, in response to real-life events. In Figure 3, we show a frequency word cloud capturing the keywords communicated by the tweets about Syria on 20 March 2013 between 8:00 pm and 10:00 pm EDT. This coincides to news breaking out of Syria reporting the suspected use of chemical weapons. In the word cloud, the bigger size words are the most popular ones, and we see how words associated with this massacre have emerged rapidly as keywords for that time period.

In our model, unique keyword instances are stored in a separate tags table, which is linked to the entries table in a many-to-many relationship. Tags and locations have a many-to-many relationship: a tag can be associated with many locations, while the same location can be associated with many tags.



Figure 3. A frequency word cloud capturing the keywords communicated by the tweets out of Syria on 20 March 2013 between 8:00 pm and 10:00 pm EDT. Keyword size reflects frequency.

3.1.6. Source component

The source component, and the corresponding sources table, serves as a directory of the different source types that can provide data to our framework. Each source can be associated with one or more feeds, which are instances of a given source type. A source feed provides content filtered according to predefined criteria. Each feed instance is characterized by a feed_id and a name. Additionally, the feed's record can also contain a query that is used to retrieve the content; however, since some APIs allow building complex queries that cannot be described by a single URL, the query field is optional.

3.2. Source-dependent components

The data received from the feeds can be separated into two parts: a part comprised of content that is common to all sources and a second part comprised of the remaining (non common) elements of the data received from the social media API source. This allows streamlining the integration of information from different sources, while maintaining the unique elements of each source. Source independent data, in the form of location, keywords, submission time, or authorship, is stored in the database in the tables described above. Source-dependent data are stored in source-specific tables. Since each of these tables is tailored to accommodate the data specific to a given source, we are not constrained by the requirement to keep uniformity between different sources in our model. For example, the Twitter API returns a tweet attribute that contains the content of a tweet, while the Flickr API returns a photo URL and has no equivalent to the tweet attribute. Such source-specific attributes, which are driven by the characteristics of each social media source, are regarded as source-dependent and are therefore stored separately. It should be noted that in the case of Flickr, our focus is on retrieving the location of images rather than retrieving them, and accordingly we do not consider the photo metadata in our architecture. Similarly, in the case of YouTube, we consider the location of videos rather than the video attributes.

Although Twitter, Flickr, and YouTube (and others) assign a unique entry identification tag any time media is fed to the service, we maintain in our model an internal unique

identification tag for each entry to ensure consistency and enable links between entries and between entries and other model components. The entry identification tags that were assigned by each social media feed are kept in conjunction to the internal unique identification tag so that the relation between our identification tags and the social media service tags are preserved.

4. Implementation and sample test cases

The prototype Geosocial Gauge system presented in this article has been implemented using a combination of PostgreSQL and MongoDB databases for data storage and programmed using Ruby, SQL, and R. The system can function using either Twitter’s free API or a commercial service to access wider data streams. When using the free API, the data stream is capped to 50 tweets per second¹⁰ and is drawn from a 1% random sample of its traffic. With Flickr there are no substantial restrictions on public content access.

In Figure 4, we show a screenshot from our system’s front-end, showing live updates of Twitter traffic. It comprises four views of the streaming data. The top left shows a heat map indicating locations of high social media activity (tweets in this case). On the top right we display streaming social media feeds, harvested by our system from the corresponding API sites. On the bottom left we show a word cloud representation of the content of these messages. It shows the 70 most frequently used words in that data corpus (word size is proportional to usage frequency). Finally, the table in bottom right shows a rank of the origin locations of these contributions. The data (top right) are streaming in continuously and so are the overlaid icons. All other three content is updated every 5 minutes using data from rolling 10-minute windows. As this frontend focuses on the dynamic nature of the data feeds, it does not show Flickr and YouTube.



Figure 4. A screenshot from the front-end of our Geosocial Gauge system. The data displayed in this figure correspond to a discussion about the unrest in Turkey and were collected on 12 June 2013.

4.1. Identifying and mapping connected communities and their structure

To demonstrate our system's capabilities, we collected over a period of approximately 3 months (1 August 2012–28 October 2012) tweets referring to a mid-Atlantic University using keywords associated with this University's name and two other characteristic nicknames (an acronym, and the University's Twitter handle). We used the free Twitter API to do so, and this resulted to a collection of 49,145 tweets. Among them, 22,693 tweets were geolocated: 854 (or 1.7% of the original data set) had precise coordinates associated with them, while 21,839 (44.38%) were geolocated using a gazetteer (Yahoo!'s PlaceFinder).

It should be noted here that the use of gazetteer for geolocation does not necessarily guarantee a foolproof result as, for example, toponyms may have multiple potential matches. The disambiguation of toponyms and the geolocation of user-contributed documents even in the absence of any such information has been the subject of substantial research, preceding the emergence of social media. Earlier work addressed the geo-tagging of web content through the use of contextual information (see, e.g., Amitay *et al.* 2004), and similar techniques have more recently emerged to address the geolocation of web blogs (see, e.g., Fink *et al.* 2009), wikipedia contributions through the use of co-occurrence models (Overell and Rüger 2008), or social media in particular through the detection of geo-scope content in tweets (Cheng *et al.* 2010). The accuracy of these techniques varies, but the potential that emerges is that such solutions can eventually be used to raise the geolocation percentage of tweets based on content rather than metadata geolocation information. Nevertheless, this remains an on-going and partially successful research issue that extends beyond the scope of this article. A comprehensive recent study, for example, assessed the rates of success of various geocoders (Roongpiboonsopit and Karimi 2010). In our case, we found that comparisons of precisely geolocated results with the corresponding gazetteer-derived products through a manual inspection allowed us to detect outliers, which are in turn manually filtered. Considering that we are focusing on geolocated communities rather than individuals, this is an approach that we consider to be adequate for our objective.

In Figure 5, we show the geographic distribution of the Twitter users who were involved in this discussion, thus mapping this online community. In our data corpus, the average number of tweet locations per user was 1.043 and the median was 1 (the maximum number of locations associated with a single user was 13). Accordingly, we argue that the tweet locations reflect the user location as well in this case. We can observe a main constituency extending primarily from Virginia to Massachusetts along the Atlantic coast and complementary distributed clusters along the Eastern Seaboard, Near Midwest, South, and California. One could argue that this is a standard human geography map, in the sense that it shows the distribution of a population that shares a common attribute (in this case, interest for a particular University). However, social media content allows us to extend this capability by incorporating connections linking these clusters, as we demonstrate below.

Figure 6 shows the structure of the social network that was mapped in Figure 5. Here nodes represent Twitter users. There exist a variety of mechanisms through which one could identify connections among these nodes, for example, through follower relations, or mentions to other users. For this study, we consider in particular retweets as suggested by Stefanidis *et al.* (2013), as they are active expressions of user connections and social networking (Boyd *et al.* 2010, Kwak *et al.* 2010, Yang *et al.* 2010). In comparison to 'follow' or '@' relations, retweets offer the advantage of reflecting an active interaction between users that is important for the dissemination of information and community building, even though it does not necessarily imply friendship or other forms of social interactions. The

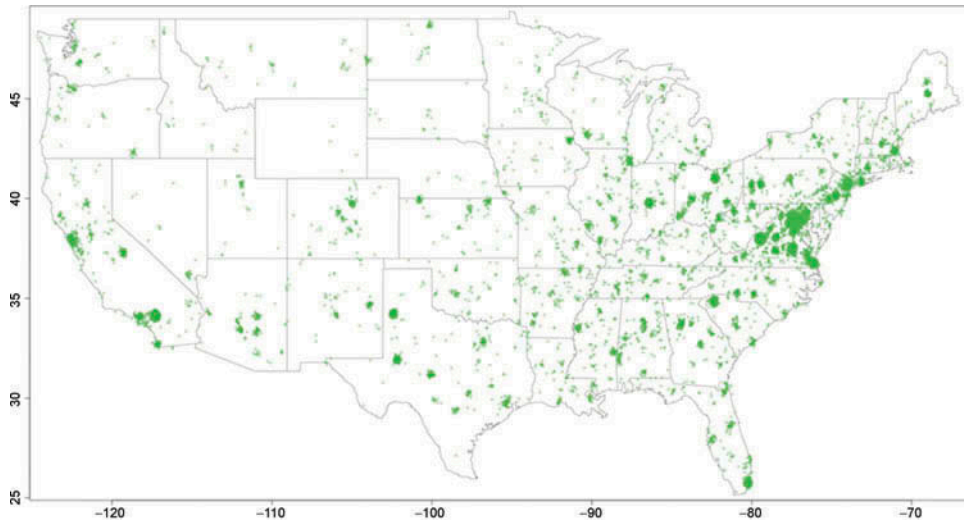


Figure 5. The point distribution of tweets in our data corpus and correspondingly a map of the involved Twitter community.

958 nodes (and corresponding users) displayed in Figure 6 are nodes of degree 3 or higher, that is, they were involved in at least 3 retweets (either retweeting or being retweeted). They represent only 17% of the total number of 5631 nodes collected in our study. The remainder 83% of nodes in our data corpus were involved in two or less retweeting activities. The connections shown in the graph of Figure 6 are 2281 or 35% of the total number of 6450 connections in our data corpus. Thus, we see that the top 17% of the nodes account for 35% of the total retweet connections. We considered 3 as the cut-off rate for this analysis to be consistent with the commonly accepted view of the minimum size of small groups and cliques (3–5 members) in social networks (e.g., Dunbar and Spoors 1995, Hill and Dunbar 2003). The number of retweets visualized (Figure 6; as connections between nodes) is 2281, out of a total of 6450 retweets. This indicates that 35% of all retweets in our data corpus were generated from the above-mentioned 17% of the nodes.

By design, the network is visualized in Figure 6 with minimally connected nodes arranged at its perimeter. In contrast, highly connected nodes, and their corresponding denser, more connected clusters, are arranged towards the center of this graph. To better communicate visually the complexity of these connections, we provide in Figure 7 a close-up of the central window identified in Figure 6. The size of node discs is proportional to their retweeting activity (either retweeting other messages or being retweeted by others). Nodes represented by bigger discs are, therefore, responsible for more information traffic within this network, either as popular contributors or as relays of information. The thickness of the connections is proportional to the volume of established links among the corresponding nodes. A pair of nodes connected through a thick line displays more retweet activity compared to other pairs connected through thin lines. In Figure 7, we observe some key nodes of our network and their connections. The two major nodes in the middle of Figure 7 correspond to the official University Twitter handle (largest node in the middle of figure) and a leading University administrator (slightly smaller node to the right of it). These two nodes share common connections, and they retweet each other's messages quite often (as indicated by the thick lines connecting them). It is also worth pointing out the pair




Figure 6. The retweet social network structure of the community. The names of the nodes have been removed to preserve anonymity.




Figure 7. A close-up of the social network of Figure 6, demonstrating the complex connections among its nodes.

of nodes at the lower right side of the graph in [Figure 7](#), where two nodes share a very strong connection as marked by the thick line.

To better understand cluster formation within such networks, we show in [Figure 8](#) a connected subset of the nodes of [Figure 6](#). Three different clusters can be recognized as




Figure 8. A small cluster of nodes showing the connections among them.

identified by different colors. They each have their own internal structure (they are typically formed around few dominant nodes, the ones with larger disc size). These communities are connected to form a larger social structure through select bridge nodes. In this figure, these bridge nodes are the two red nodes that connect the red cluster to the green and yellow ones. It is worth mentioning here that these clusters are identified based solely on manifested connections among their nodes (in the form of retweets), regardless of a specific discussion sub-topic. The graph of [Figure 8](#) was built using a depth search approach (Tarjan 1972), but one could use a variety of approaches to do so. Duch and Arenas (2005) provide a good overview of techniques for connected community detection in social networks.

Once such network structure has been extracted, we can proceed with augmenting the simple map of [Figure 5](#) by introducing connections among nodes to reveal *links between spaces*, as we show in [Figure 9](#). The connecting lines in this Figure are visualizations in space of the connections of the social network of [Figure 6](#). Through these links, different locations are linked by virtue of the activities undertaken by their occupants.

We see quite vividly how the epicenter of this community is in the Capital area, and the satellite communities are connected to it. We can also observe that social proximity often overtakes geographical proximity when such connections are established. In an earlier article, Takhteyev *et al.* (2012) had shown that when it comes to Twitter communities, a substantial share of ties lies within the same metropolitan region, and that between regional clusters, distance, national borders and language differences all predict Twitter ties. Our data show that thematic affinity deserves at least an equal level of consideration as these other parameters. For example, we observe that the Los Angeles metro area community

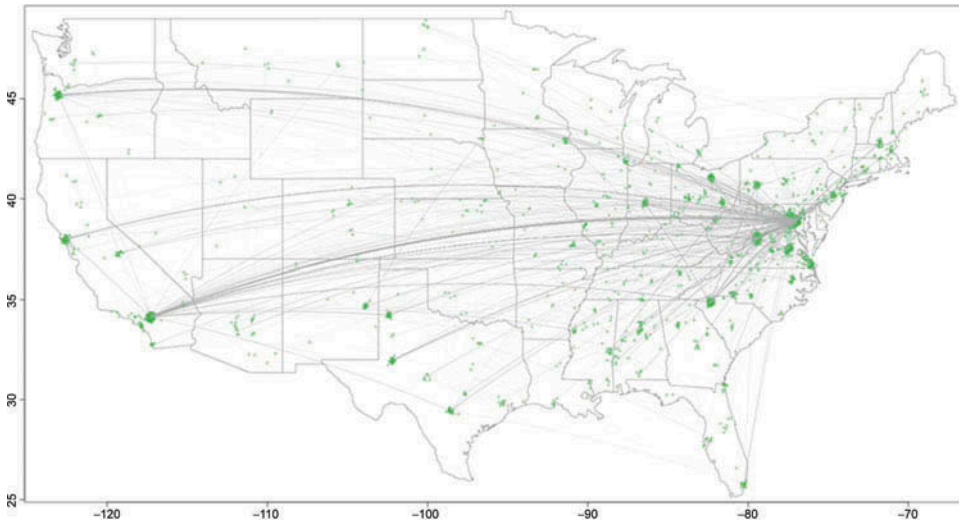


Figure 9. Mapping the connections among the nodes of Figure 6 to reveal the underlying links among the corresponding locations.

is more closely connected to the central community in the DC metro area, rather than to the San Francisco metro area community (as indicated by the thickness of lines connecting these communities). Even though both California communities share an interest on the same topic, the relative significance of the DC community (being local to the University of interest) overtakes in this particular case the parameters presented by Takhteyev *et al.* (2012) to reshape the geography of this particular Twitter network. This results in two large communities in California that are minimally connected directly to each other. Instead, they are connected rather indirectly, through the central community in the Capital area.

We should mention here that the green nodes of Figure 9 are a subset of the ones in Figure 6, as they only show pairs where *both* originator and retweeter were geolocated. This accounts for 2552 retweets (among the total of 6450 retweets). In contrast, the social network structure of Figure 6 included both geolocated and non-geolocated nodes. Thicker lines in Figure 9 indicate multiple connections among a pair of locations (and corresponding nodes associated with these locations). This visualization provides a novel view of space, as linked locations, and represents an advancement compared to our traditional, strict Cartesian perception of it. It reflects the transformation of a simple map of individual users (as shown in Figure 5) to a map of a connected community and its links and ushers a new, geosocial view of space. One can argue that this is in essence the snapshot of a *geosocial ecosystem*, comprising geolocated nodes (reflecting individual members) and connections expressing how these nodes communicate with each other.

4.2. Events and cross-platform knowledge discovery

The geosocial system detected and mapped in Section 4.1 comprises a core component in the form of a basic community that forms it, its interests as they are expressed through the associated discussion keywords, and the links that are established through communications within that community. This core component of our geosocial system is occasionally disturbed by events. Here we will attempt to highlight how such events impact our system through a representative example.




Figure 10. Twitter traffic associated with our study (green line) and the portion of these tweets that include a reference to President Obama (blue line). The x -axis represents time (from 5 August 2012 to 10 October 2012), and the y -axis represents the tweet count. The two top traffic dates in our data corpus are marked by red circles.

In Figure 10, we show Twitter traffic associated with our system (green line). We can easily recognize two peaks, marked by circles, at 5 October 2012 and 19 October 2012. The keyword analysis of the data reveals that a new keyword has emerged on these dates (namely ‘Obama’), and the blue line in the same graph shows the portion of traffic that is associated with this new keyword. We can see that while the President’s name was minimally involved with information traffic in our geosocial system before these dates, it accounts for approximately a third of the corresponding traffic on 5 October and 19 October (growing to that level over a period of two days before each peak).

These results produced a new set of keywords, namely the university name and the word ‘Obama’ that were used for a subsequent social media query as part of the knowledge discovery process within GeoSocial Gauge. In this case, the Flickr API was called using these new keywords and a temporal range corresponding to the period of interest based on the analysis of Twitter results. In this particular case, limiting the search to the peak dates we find in Flickr clusters of images contributed on these dates. For example, in the second date, the GeoSocial Gauge harvested 49 geolocated images from Flickr. After plotting their locations (in Google Earth), we see that they are spatially clustered at a football field of the University under observation as shown in Figure 11. Querying the YouTube API with the same parameters at the same time did not produce any geolocated videos. A recent study has indicated that approximately 4.5% of Flickr and 3% of YouTube content is geolocated (Friedland and Sommer 2010).

An online news search shows that on both days President Obama addressed election rallies at this University Campus. On the second day, he gave a speech to an audience of few thousands at this particular practice field as shown in Figure 12. Thus, we observe that Twitter traffic variations are good indicators for event detection (e.g., by identifying spikes and corresponding emerging keywords), while relevant Flickr clusters can support the discovery of the location of these events.

To examine the effect of this exogenous event onto the geosocial system that we investigated, we plotted in Figure 13 our social network but marked in blue the nodes that only contributed tweets related to this particular event (i.e., all University-related tweets

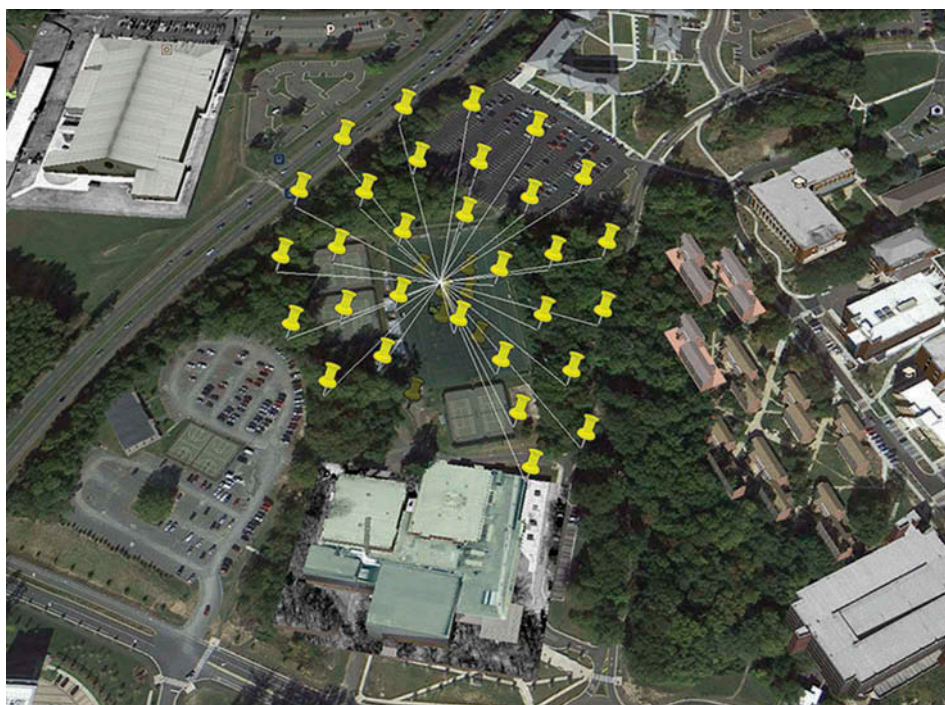


Figure 11. Clusters of geolocated images contributed on 19 October 2012. We can identify images contributed from eight different locations, including a large cluster of 30 images contributed from a single user at a single location.



Figure 12. The election rally event that was documented in Flickr and gave rise to the traffic peak of 19 October 2012. It is interesting to observe the numerous handheld devices used to document the event.




Figure 13. Another layout of the network of Figure 6. Here we identify by blue the nodes (and retweets) that have only contributed tweets related to the exogenous event of October 19, 2012.

contributed by them included the President's name). We observe that these blue nodes represent a very small portion of our network, and they are minimally connected within it, with the exception of a small cluster of 4 blue nodes within the marked box of Figure 13. Only 18 blue nodes are present in Figure 13, that is, they have at least 3 connections. The total number of blue nodes in our data corpus was 315, leaving 297 blue nodes with 2 or less connections. Overall, the blue nodes are sparsely connected, with only 148 connections among them, as shown in Figure 14. This simple example highlights a process through which it is possible to monitor, study, and assess the impact of exogenous events on a geosocial ecosystem. In our particular test case, the event affected traffic patterns by introducing spikes as shown in Figure 10, but had little effect on the actual overall structure of our social network.

5. Conclusion and outlook

More than 2300 years ago, Aristotle argued that 'Man is by nature a social animal; an individual who is unsocial naturally and not accidentally is either beneath our notice or more than human' (Vinciarelli *et al.* 2012), positioning social interactions as one of our most basic needs as individuals and as a society. Today we continue to engage in social interactions, and our beliefs, behavior, feelings, and actions are still deeply influenced by the people we interact with, whether they are our family, our friends, our peers, or society at large. However, the medium through which we can engage others has changed dramatically: in addition to physical interaction, social media in its many forms has opened an entirely new avenue to enable an alternative form of interaction. A recent survey on the use of social media among adult US users (Smith 2011) indicated that staying in touch with




Figure 14. The sparse connections among the 315 blue nodes. A comparison to Figure 6 shows the sparsity of this sub-network.

current friends (67%), family members (64%), and long lost old friends (50%) were the three most significant reasons for using social media. The uniqueness of social media lies not only in its ability to foster further social interaction, but also in its potential to transform the way we study it: for the first time, we are able to closely *observe* Aristotle's social animal and *measure* and *quantify* interactions and behaviors. Furthermore, the widespread use of social media and its ubiquity in virtually every aspect of modern society are making possible, for the first time, to gain valuable insights into the links and interactions between people in a variety of settings and contexts and the *physical space* surrounding them. Thus, information harvesting from social media emerges as a new instrument for observing the human landscape, the geosocial ecosystems that are formed, and their evolution over time, at a larger scale and at a finer resolution than ever before.

In this article, we presented how such an instrument – Geosocial Gauge – can be built and utilized to derive geospatial knowledge from social media feeds. In particular, we addressed the challenge of harvesting social media feeds and how such data can be processed and analyzed to reveal information about communities formed through them, their social networks, and the geographical space in which they operate. In addition, we presented a conceptual model for storing and managing heterogeneous geosocial data in a unified database and provided a detailed description of its various source-dependent and source-independent components. This approach allows bridging the gap between independent spatial (e.g., simple map mash-ups) and social network (e.g., Brabasi 2002) analyses and supports the discovery of additional knowledge. To showcase how Geosocial Gauge can be used to glean geosocial knowledge, we presented two sample test cases – identifying and mapping connected communities around an academic institution and discovering and locating events affecting such a community, using diverse social media sources. This

allowed us to gain knowledge about the way in which these community members are distributed in space. This is followed by a study of their social network structure, which is then projected onto space to discover knowledge on the ties among distributed clusters of this community. By gleaning such knowledge, we are then able to see how this community is affected by relevant events. In doing so, we gain an understanding of how humans act as geosensors to report on a broad array of sociocultural events affecting them. This understanding of the connections between space, social networks, and events will help our community extract knowledge from the unstructured data of social media feeds.

The unprecedented ability to harvest, process, and derive knowledge from geosocial media is challenging the limits of our capacity to understand the geosocial realm. While in the past, human interaction occurred only in the physical space, today much of it also occurs in virtual space. The interplay between these two spaces is still not well understood, calling for new theories for grounding these phenomena and their relation to space, human behavior, and social interaction to link the virtual and the physical domains. In addition, the sheer volume of geosocial media puts forward some significant challenges in the development of automated tools for processing and deriving knowledge from such feeds.

As we proceed exploring these opportunities, it is important that we remain cognizant of the associated privacy issues to ensure its proper use. The challenge exceeds beyond the simple anonymization of such data. Studies have shown that by exploiting links in social networks, public profiles can be exploited to discover hidden private attributes in social media platforms (Zheleva and Getoor 2009). The privacy concerns are further complicated by the presence of geotagged information at-large (e.g., Friedland and Sommer 2010). For example, recent studies have shown that the analysis of human mobility data in the form of cellphone usage allows for the unique identification of individuals by using as few as four spatiotemporal points in these trajectories, even when coarse geolocation information is made available (de Montjoy *et al.* 2013). Accordingly, proposed solutions for privacy-aware collection of aggregate spatial data (e.g., Xie *et al.* 2011) are of questionable effectiveness when it comes to social media content. The broad range of information that is communicated through social media, an aggregate of location, social connections, and personal views, is accentuating the need to re-conceptualize the concept of privacy, as suggested by Elwood and Leszczynski (2011) with respect to the geoweb.

While data harvested from social media holds great potential for unveiling valuable geosocial knowledge, it is important to recognize that the population of social media users may be demographically skewed. For example, a recent survey in the United States showed that Twitter is especially appealing to adults between the ages of 18 to 29, African-Americans, and urban residents, while Instagram may be particularly appealing also to Hispanic and women population in this age group (Duggan and Brenner 2012). A similar study indicated that over 65% of its user constituency being younger than 44 years and that usage is notably higher in urban areas (Smith and Brenner 2012). Similarly, Mislove *et al.* (2011) found that in the United States, densely population regions are overrepresented in Twitter and that its users are predominantly male. The effect of such demographic biases can be significant, as was recently shown in the context of electoral predictions using social media (e.g., Gayo-Avello *et al.* 2011). However, in other cases it was found that the sheer volume of social media data could lead to results of similar predictive power as traditional election polls (e.g., Tumasjan *et al.* 2011). As our community is proceeding with the further study of the content of social media contributions to derive valuable geospatial knowledge, we anticipate a more in-depth analysis of the various normalization aspects of such data.

With millions of tweets and thousands of Flickr images posted hourly, social media content is rapidly emerging as a new big data challenge for the computational and geospatial

communities. While this work marks a step in this direction, further work is required with respect to managing and analyzing social media feeds. This is further emphasized by the ever-evolving nature and complexity of social media feeds. For example, comments within YouTube or Flickr may also be mined to further explore social interactions among users (e.g., Siersdorfer *et al.* 2010), while Instagram emerged in the past year a new complementary source of imagery (Hochman and Schwartz 2012). It is also important to note that as the human landscape can be highly dynamic, especially in the case of disruptive events (e.g., an earthquake or a geopolitical conflict), the dynamic nature of geosocial information and the knowledge that is derived from it should also be addressed. This can include, for example, studying and developing new tools for measuring and monitoring the resilience of ties in the social and geographic domains and how links in these two domains relate to one another over time and space. Combined, these emerging opportunities promise to substantially advance the study of the human landscape at spatial and temporal scales and resolutions unfathomed so far.

Notes

1. <http://www.ushahidi.com>
2. <http://140kit.com/>
3. <http://hootsuite.com/>
4. <http://dbpedia.org/>
5. <http://www.freebase.com/>
6. <http://www.geonames.org/>
7. <http://about.topsy.com/technology/overview/>
8. <http://www.metacarta.com/products-platform-geotag.htm>
9. <http://bit.ly/S01Vlr>
10. This information is current and is subject to change by Twitter itself.

References

- Ames, M. and Naaman, M., 2007. Why we tag: motivations for annotation in mobile and online media. In: *Proceedings of ACM SIGCHI conference on human factors in computing systems*, San Jose, CA, 971–980.
- Amitay, E., *et al.*, 2004. Web-a-where: geotagging web content. In: *Proceedings of the 27th annual international ACM SIGIR conference on research and development in information retrieval*, Sheffield, UK, 273–280.
- Backstrom, L., Sun, E., and Marlow, C., 2010. Find me if you can: improving geographical prediction with social and spatial proximity. In: *Proceedings of the 19th international conference on world wide web*, Raleigh, NC, 61–70.
- Barabasi, A., 2002. *Linked: the new science of networks*. New York, NY: Perseus Publishing.
- Blasingame, D., 2011. Twitter first: changing TV news 140 characters at a time. In: *Proceedings of the 12th international symposium on online journalism*, Austin, TX.
- Bonchi, F., *et al.*, 2011. Social network analysis and mining for business applications. *ACM Transactions on Intelligent Systems and Technology*, 2 (3), doi:10.1145/1961189.1961194
- Boyd, D., Golder, S., and Lotan, G., 2010. Tweet, tweet, retweet: conversational aspects of retweeting on Twitter. In: *Proceedings of the 43rd IEEE Hawaii international conference on system sciences*, Kauai, HI.
- Bruns, A. and Burgess, J.E., 2011. #Ausvotes: how Twitter covered the 2010 Australian federal election. *Communication, Politics and Culture*, 44 (2), 37–56.
- Cheng, Z., Caverlee, J., and Lee, K., 2010. You are where you tweet: a content-based approach to geolocating Twitter users. In: *Proceedings of the ACM conference on information and knowledge management*, Toronto, Canada, 759–768.
- Croitoru, A., *et al.*, 2012. *Towards a collaborative geosocial analysis workbench*. Washington, DC: COM-Geo.

- Crooks, A.T., et al., 2013. #Earthquake: Twitter as a distributed sensor system. *Transactions in GIS*, 17 (1), 124–147.
- deMontjoye, Y. A., et al., 2013. Unique in the crowd: the privacy bounds of human mobility. *Scientific Reports*, 3 (Article No. 1376).
- Duch, J. and Arenas, A., 2005. Community detection in complex networks using extremal optimization. *Physical Review E*, 72 (2), 027104.
- Duggan, M. and Brenner, J., 2012. *The demographics of social media users – 2012*. Pew Research Center, Washington, DC. Available from: <http://bit.ly/XORHo0> [Accessed 19 January 2013].
- Dunbar, R.I.M. and Spoor, M., 1995. Social networks, support cliques, and kinship. *Human Nature*, 6 (3), 273–290.
- Earle, P., et al., 2010. OMG earthquake! Can Twitter improve earthquake response? *Seismological Research Letters*, 81 (2), 246–251.
- Earle, P., Bowden, D.C., and Guy, M., 2011. Twitter earthquake detection: earthquake monitoring in a social world. *Annals of Geophysics*, 54 (6), doi:10.4401/ag-5364
- Elwood, S. and Leszczynski, A., 2011. Privacy, reconsidered: new representations, data practices, and the geoweb. *Geoforum*, 42 (1), 6–15.
- Fielding, R.T., 2000. *Architectural styles and the design of network-based software architectures*. Thesis (PhD). University of California, Irvine, CA.
- Fink, C., et al., 2009. The geolocation of web logs from textual clues. In: *Proceedings international conference on computational science and engineering*, Vancouver, Canada, 1088–1092.
- Forbes, 2012. Twitter's Dick Costolo: Twitter mobile ad revenue beats desktop on some days. Available from: <http://onforb.es/KgTWYP> [Accessed 19 January 2013].
- Friedland, G. and Sommer, R., 2010. Cybercasing the joint: on the privacy implications of geotagging. In: *Proceedings of the fifth USENIX workshop on hot topics in security (HotSec 10)*, Washington, DC.
- Fritz, S., et al., 2009. Geo-Wiki.Org: the use of crowdsourcing to improve global land cover. *Remote Sensing*, 1 (3), 345–354.
- Gayo-Avello, D., Metaxas, T., and Mustafaraj, E., 2011. Limits of electoral predictions using Twitter. In: *Proceedings of the 5th international AAAI conference on weblogs and social media (ICWSM)* Barcelona, Spain, AAA Press, 490–493.
- Goodchild, M.F., 2007. Citizens as sensors: the world of volunteered geography. *GeoJournal*, 69 (4), 211–221.
- Goodchild, M.F. and Glennon, J.A., 2010. Crowdsourcing geographic information for disaster response: a research frontier. *International Journal of Digital Earth*, 3 (3), 231–241.
- Hecht, B., et al., 2011. Tweets from Justin Bieber's heart: the dynamics of the 'location' field in user profiles. In: *Proceedings of the ACM CHI conference on human factors in computing systems*, Vancouver, Canada, 237–246.
- Hill, R.A. and Dunbar, R.I.M., 2003. Social network size in humans. *Human Nature*, 14 (1), 53–72.
- Hochman, N. and Schwartz, R., 2012. Visualizing instagram: tracing cultural visual rhythms. In: *Proceedings of the sixth international AAAI conference on weblogs and social media*, Dublin, Ireland, 6–9.
- Hollis, C., 2011. *2011 IDC digital universe study: big data is here, now what?* Available from: <http://bit.ly/kouTgc> [Accessed 19 January, 2013].
- Huang, J., Thornton, K., and Efthimiadis, E., 2010. Conversational tagging in Twitter. In: *Proceedings of the 21st ACM conference on hypertext and hypermedia*, Toronto, Canada, 173–178.
- Iofciu, T., et al., 2011. Identifying users across social tagging systems. In: *Proceedings of the 5th international AAAI conference on weblogs and social media*, Barcelona, Spain, 522–525.
- Java, A., et al., 2009. Why we Twitter: an analysis of a microblogging community. In: R. Goebel, J. Siekmann, and W. Wahlster, eds. *Advances in web mining and web usage analysis*, Lecture Notes in Computer Science. Vol. 5439. Berlin, Germany: Springer, 118–138.
- Kaigo, M., 2012. Social media usage during disasters and social capital: Twitter and the great East Japan earthquake. *Keio Communication Review*, 34, 19–35.
- Kalnikaite, V., et al., 2010. Now let me see where I was: understanding how lifelogs mediate memory. In: *Proceedings of the 28th international conference on human factors in computing systems*, Atlanta, GA, 2045–2054.
- Kelley, M.J., 2011. The emergent urban imaginaries of geosocial media. *GeoJournal*, 78 (1), 181–203.

- Kraut, R.E., et al., 2013. *Public response to alerts and warnings using social media: report of a workshop on current knowledge and research gaps*. Washington, DC: National Academies Press.
- Kwak, H., et al., 2010. What is Twitter, a social network or a news media? In: *Proceedings of the 19th international conference on world wide web*, Raleigh, NC, 591–600.
- Larson, M., Soleymani, M., and Serdykov, P., 2011. Automatic tagging and geotagging in video collections and communities. In: *Proceedings of the 1st ACM international conference on multimedia retrieval*, Trento, Italy.
- Leetaru, K., et al., 2013. Mapping the global twitter heartbeat: the geography of Twitter. *First Monday*, 18 (5), doi:10.5210/fm.v18i5.4366
- Lieberman, M.D., Samet, H., and Sankaranarayanan, J., 2010. Geotagging with local lexicons to build indexes for textually-specified spatial data. In: *Proceedings of the IEEE 26th international conference on data engineering*, Long Beach, CA, 201–212.
- Lyons, K. and Lessard, L., 2012. S-FIT: a technique for integrating social features in existing information systems. In: *Proceedings of the 2012 iConference*, New York, NY, 263–270.
- MacEachren, A.M., et al., 2011a. Senseplace2: Geotwitter analytics support for situational awareness. In: S. Miksch and M. Ward, eds. *IEEE conference in visual analytics science and technology (VAST)*, Providence, RI, 181–190.
- MacEachren, A.M., et al., 2011b. Geo-Twitter analytics: applications in crisis management. In: *Proceedings of the 25th international cartographic conference*, Paris, France, 3–8 July 2011.
- Marcus, A., et al., 2011. Processing and visualizing the data in Tweets. *SIGMOD Record*, 40 (4), 21–27.
- Mearian, L., 2011. World's data will grow by 50x in next decade, IDC study predicts, *Computer World Magazine*. Available from: <http://bit.ly/k1Jo0V>.
- Mislove, A., et al., 2011. Understanding the demographics of Twitter users. In: *Proceedings of the 5th international AAAI conference on weblogs and social media (ICWSM)*, Barcelona, Spain: AAA Press, 554–557.
- New York Times, 2011. Spotlight again falls on web tools and change, *New York Times*, Available from: <http://nyti.ms/hWnxSp2> [Accessed 23 July 2012].
- Overell, S. and Rüger, S., 2008. Using co-occurrence models for placename disambiguation. *International Journal of Geographical Information Science*, 22 (3), 265–287.
- Pollock, J., 2011. *Streetbook: how Egyptian and Tunisian youth hacked the Arab spring*. Technology Review. Available from: <http://bit.ly/nrO9PU>.
- Rahm, E. and Do, H.H., 2000. Data cleaning: problems and current approaches. *IEEE Data Engineering Bulletin*, 24 (4), 3–13.
- Reinhardt, W., et al., 2010. Modeling, obtaining and storing data from social media tools with artefact-actor-networks. In: *Proceedings of ABIS 2010: the 18th international workshop on personalization and recommendation on the web and beyond*, Kassel, Germany.
- Romero, D., Meeder, B., and Kleinberg, J., 2011. Differences in the mechanics of information diffusion across topics: idioms, political hashtags, and complex contagion on Twitter. In: *Proceedings of the 20th international conference on world wide web*, Hyderabad, India, 695–704. doi:10.1145/1963405.1963503
- Roongpiboonsopit, D. and Karimi, H.A., 2010. Comparative evaluation and analysis of online geocoding services. *International Journal of Geographical Information Science*, 24 (7), 1081–1100.
- Sahito, F., Latif, A., and Slany, W., 2011. Weaving Twitter stream into linked data: a proof of concept framework. In: *Proceedings of the 7th international conference on emerging technologies*, Islamabad, Pakistan, 1–6. doi:10.1109/ICET.2011.6048497.
- Sapiro, G., 2011. Images everywhere: looking for models: technical perspective. *Communications of the ACM*, 54 (5), 108–108.
- Shimojo, A., Kamada, S., Matsumoto, S., and Nakamura, M., 2010. On integrating heterogeneous lifelog services. In: *Proceedings of the 12th International Conference on Information Integration and Web-based Applications (iiWAS '10)*. ACM, New York, NY, USA, 263–272. doi:10.1145/1967486.1967529
- Siersdorfer, S., Chelaru, S., Nejdil, W., and San Pedro, J., 2010. How useful are your comments?: analyzing and predicting youtube comments and comment ratings. In: *Proceedings of the 19th international Conference on World Wide Web*, Raleigh, NC, 891–900.

- Smith, A., 2011. *Why Americans use social media: social networking sites are appealing as a way to maintain contact with close ties and reconnect with old friends*. Pew Research Center, Washington DC. Available from: <http://bit.ly/rLCsA6> [Accessed 20 January 2013].
- Smith, A. and Brenner, J., 2012. *Twitter use 2012*. Pew Research Center, Washington DC. Available from: <http://bit.ly/JTpHsO> [Accessed 20 January 2013].
- Stefanidis, T., Crooks, A.T., and Radzikowski, J., 2013. Harvesting ambient geospatial information from social media feeds. *GeoJournal*, 78 (2), 319–338.
- Sui, D. and Goodchild, M.F., 2011. The convergence of GIS and social media: challenges for GIScience. *International Journal of Geographical Information Science*, 25 (11), 1737–1748.
- Sun, A., et al., 2011. Tag-based social image retrieval: an empirical evaluation. *Journal of the American Society for Information Science and Technology*, 62 (12), 2364–2381.
- Takhteyev, Y., Gruzd, A., and Wellman, B., 2012. Geography of Twitter networks. *Social Networks*, 34 (1), 73–81.
- Tan, S., et al., 2011. Using rich social media information for music recommendation via hypergraph model. *ACM Transactions on Multimedia Computing, Communications, and Applications*, 7S (1), doi:10.1145/2037676.2037679
- Tarjan, R., 1972. Depth-first search and linear graph algorithms. *SIAM Journal on Computing*, 1 (2), 146–160.
- Tomaszewski, B., et al., 2011. Supporting geographically-aware web document foraging and sensemaking. *Computers, Environment and Urban Systems*, 35 (3), 192–207.
- Tumasjan, A., et al., 2011. Election forecasts with Twitter: how 140 characters reflect the political landscape. *Social Science Computer Review*, 29 (4), 402–418.
- Twitter, 2011. *2000 million Tweets per day*. Available from: <http://bit.ly/laY1Jx>.
- Valkanias, G. and Gunopulos, D., 2013. A UI prototype for emotion-based event detection in the live web. In: *Proceedings of the 2013 IEEE international conference on human factors in computing and informatics, special session on human-computer interaction & knowledge discovery*, Maribor, Slovenia.
- Vinciarelli, A., et al., 2012. Bridging the gap between social animal and unsocial machine: a survey of social signal processing. *IEEE Transactions on Affective Computing*, 3 (1), 69–87.
- Xie, H., Kulik, L., and Tanin, E., 2011. Privacy-aware collection of aggregate spatial data. *Data & Knowledge Engineering*, 70 (6), 576–595.
- Yang, Z., et al., 2010. Understanding retweeting behaviors in social networks. In: *Proceedings of the 19th ACM international conference on information and knowledge management*, Toronto, Canada, 1633–1636.
- YouTube, 2013. *YouTube pressroom statistics*. Available from: <http://bit.ly/gzYBVx> [Accessed 20 January 2013].
- Zhao, S., et al., 2011. *Human as real-time sensors of social and physical events: a case study of twitter and sports games*. Technical Report TR0620-2011, Houston, TX: Rice University and Motorola Labs, arXiv:1106.4300v1.
- Zheleva, E. and Getoor, L., 2009. To join or not to join: the illusion of privacy in social networks with mixed public and private user profiles. In: *Proceedings of the 18th international conference on world wide web*, Madrid, Spain, 531–540.